

Fundamentos de Machine Learning para Geometalurgia **Regression Models**

24 de abril al 4 de mayo 2023

Agenda

Machine Learning basis



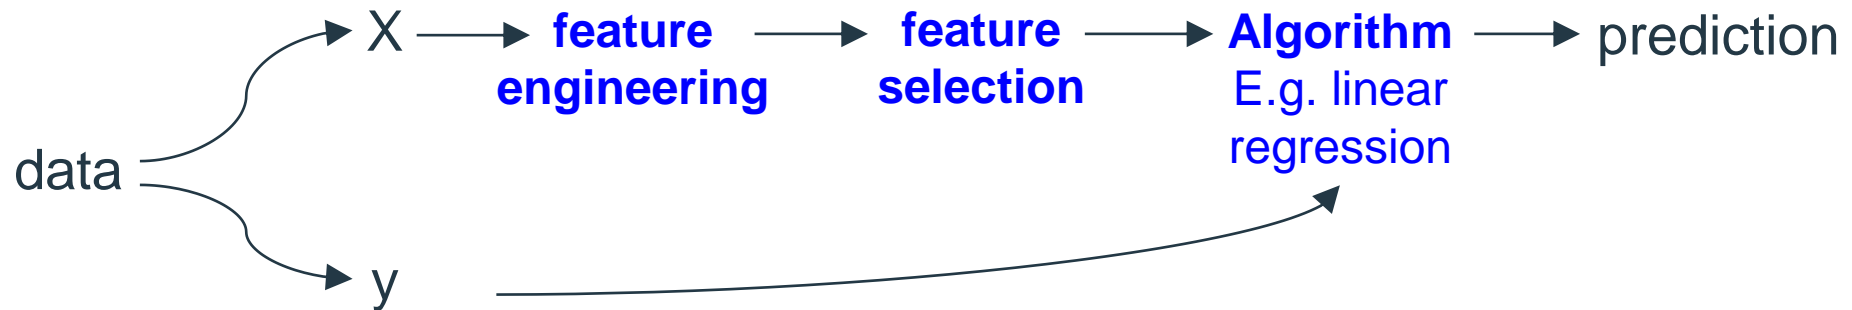
Case study

**Univariate
Exploratory Data Analysis (EDA)**

Data Preparation

**Regression model (proxy) for
geometallurgical parameter A_i**

Machine Learning Process



X
(predictor matrix)

information such as
mineralogy,
minzone, etc.

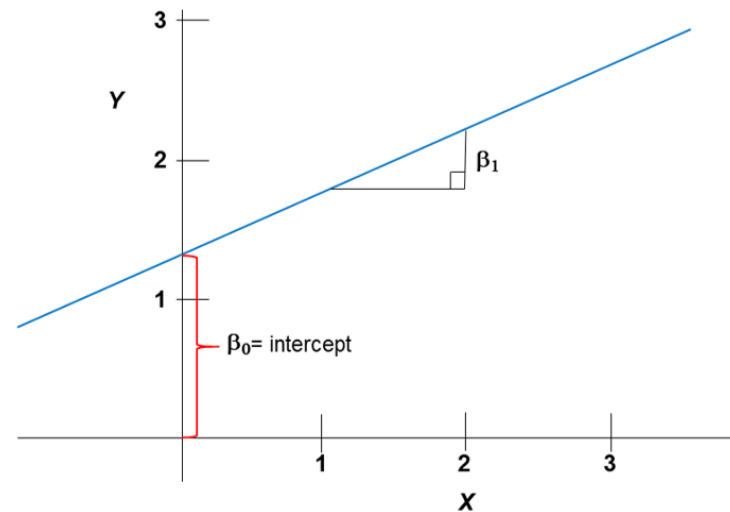
y
(target variable)

A_i values

Algorithm: Linear regression

Simple linear regression can be expressed as a function of slope (β_1) and the intercept (β_0) of a straight line:

$$Y = \beta_0 + \beta_1 X$$



Algorithm: Linear regression

Multiple linear regression model with response Y and terms X_1, \dots, X_p can be expressed as:

$$y = \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \beta_0$$

where

y : response variable

n : number of features

x_n : n -th feature

β_n : regression coefficient (weight) of the n -th feature











β_0 : y-intercept

Animation:

https://aegis4048.github.io/mutiple_linear_regression_and_visualization_in_python

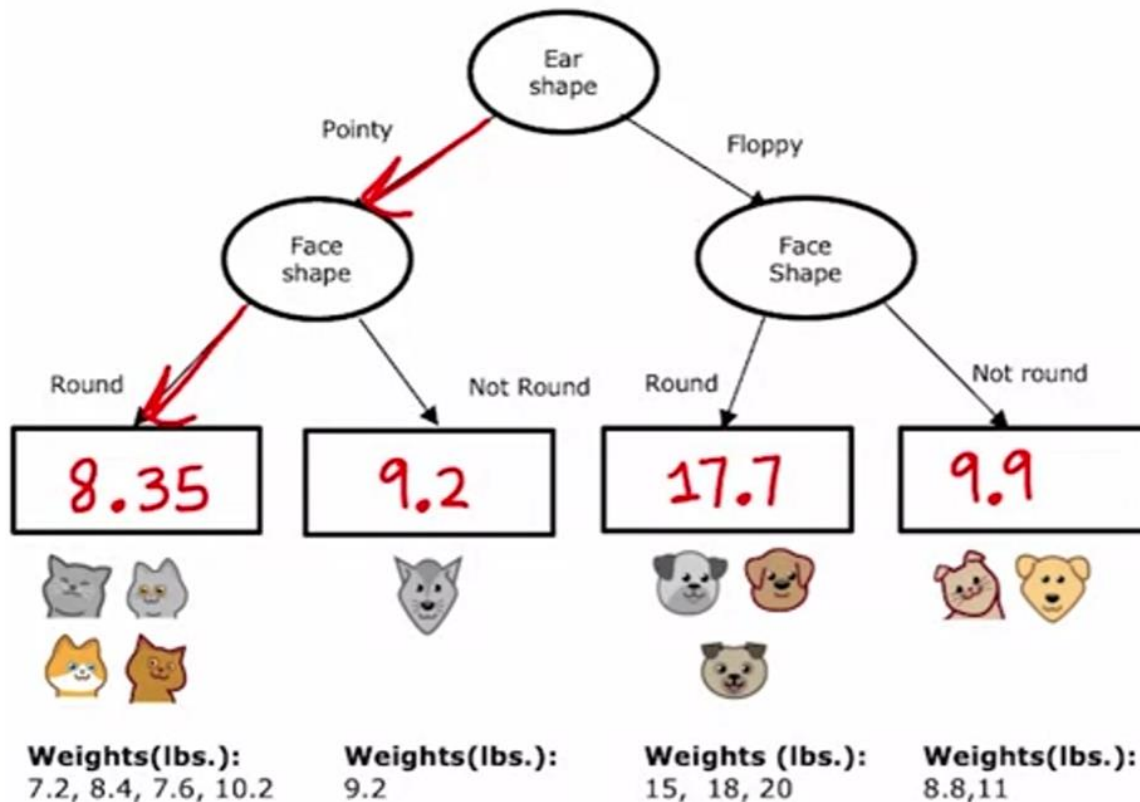
Algorithm: regression tree

Regression with Decision Trees: Predicting a number

| | Ear shape | Face shape | Whiskers | Weight (lbs.) |
|---|-----------|------------|----------|---------------|
|  | Pointy | Round | Present | 7.2 |
|  | Floppy | Not round | Present | 8.8 |
|  | Floppy | Round | Absent | 15 |
|  | Pointy | Not round | Present | 9.2 |
|  | Pointy | Round | Present | 8.4 |
|  | Pointy | Round | Absent | 7.6 |
|  | Floppy | Not round | Absent | 11 |
|  | Pointy | Round | Absent | 10.2 |
|  | Floppy | Round | Absent | 18 |
|  | Floppy | Round | Absent | 20 |

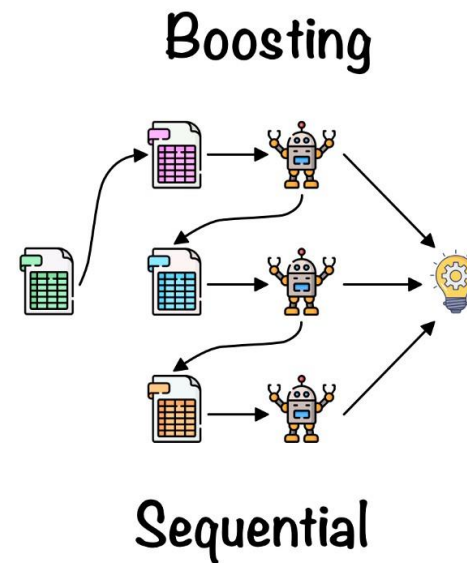
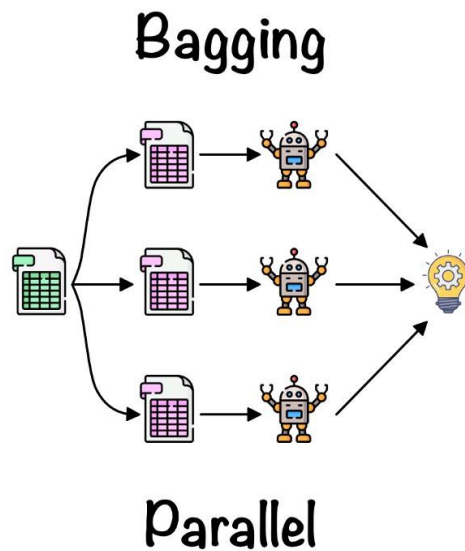
Algorithm: regression tree

Built through binary recursive partitioning, and then continues splitting each partition into smaller groups.



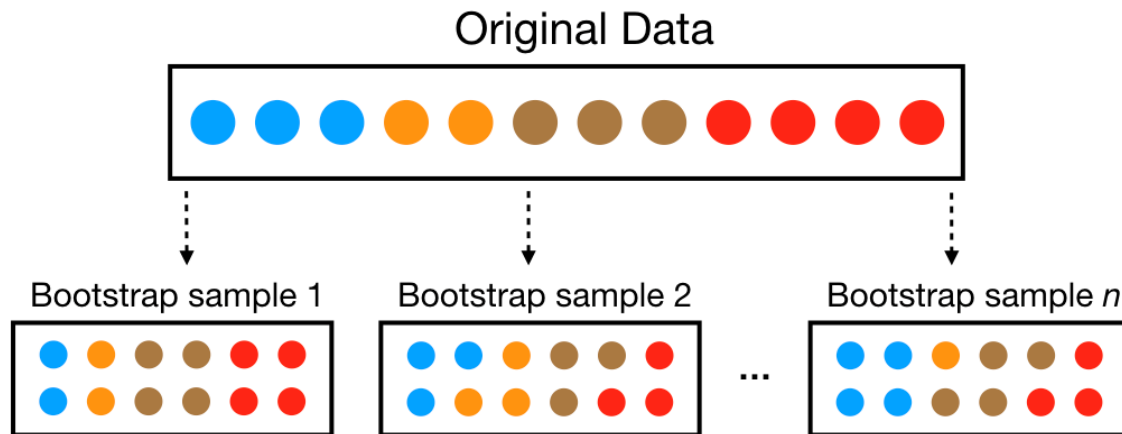
Algorithm: random forest

Ensemble learning is the process of using multiple models, trained over the same data, averaging the results to find a better predictive result. Combining weak learners to build a stronger learner usually increase the model performance. Random forest is a bagging ensemble.



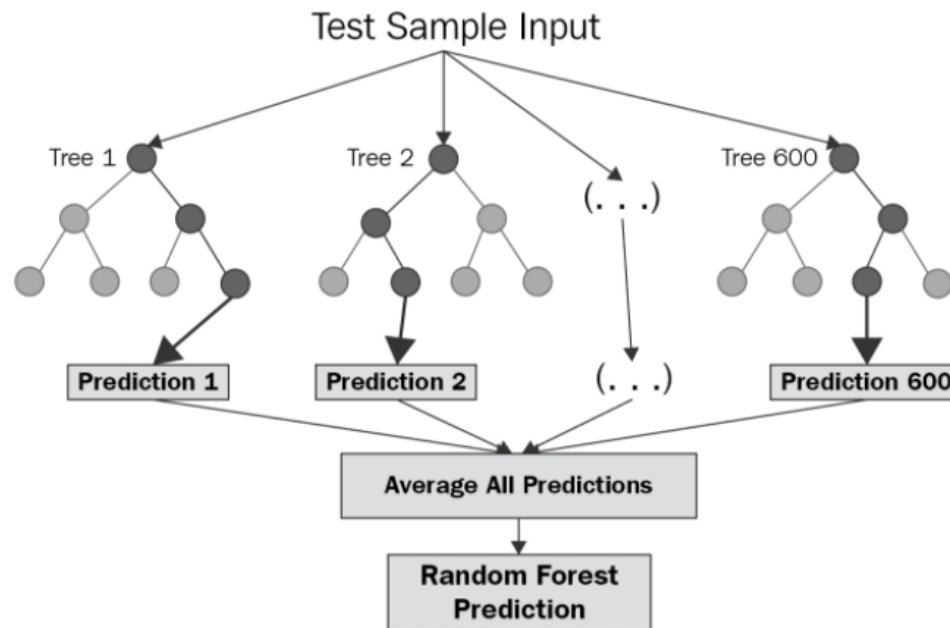
Algorithm: random forest

Bootstrapping is the process of randomly sampling subsets of a dataset over a given number of iterations and a given number of variables. Since samples are drawn with replacement, each bootstrap sample is likely to contain duplicate values.

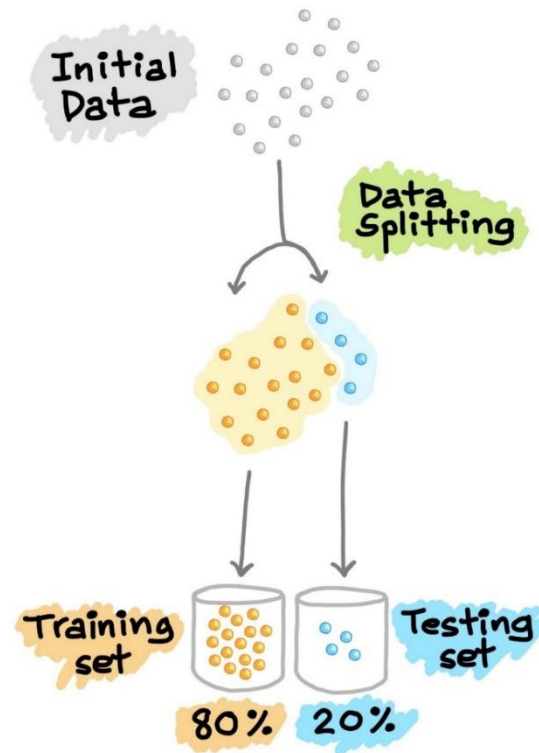


Algorithm: random forest

Bootstrapping algorithm that ensemble multiple randomly drawn decision trees from the data, averaging the results to obtain the prediction. In addition, a subset of the features is randomly selected at each node.



Data splitting



| | Features | Labels |
|--------------|----------|---------|
| Training Set | X_train | y_train |
| Test Set | X_test | y_test |

Trained model must perform well on new, unseen data. In order to simulate the new, unseen data, the available data is subjected to data splitting whereby it is split into 2 portions. 80% of the original data is used as the training set and the remaining 20% is used as the testing set

Tune hyperparameters with GridsearchCV

Hyperparameters are variables that the user specify usually while building the Machine Learning model and are used to evaluate optimal parameters of the model. Example: **max_depth** in Random Forest. But, How can we find the best hyperparameters values to get the best prediction results from our model?

Grid Search uses a different combination of all the specified hyperparameters and their values and calculates the performance for each combination and selects the best value for the hyperparameters.

Tune hyperparameters with GridsearchCV

In GridSearchCV, along with Grid Search, cross-validation (CV) is also performed. In CV, train data is divided into two parts: train data and validation (test) data.

GridsearchCV

```
# Instantiate the grid search model
rf = RandomForestRegressor(random_state=1)

param_grid = {
    'max_depth': [10, 20, 30, 40],
    'max_features': [3, 4, 5, 6, 7],
    'min_samples_leaf': [2, 3, 4, 5],
    'min_samples_split': [3, 4, 5, 6],
    'n_estimators': [20, 30, 40, 50]
}

grid = GridSearchCV(estimator=rf,
                    param_grid=param_grid,
                    cv=5, verbose=1)
```

hyperparameters

**number of groups or folds
for cross-validation**

Cross-validation (CV)

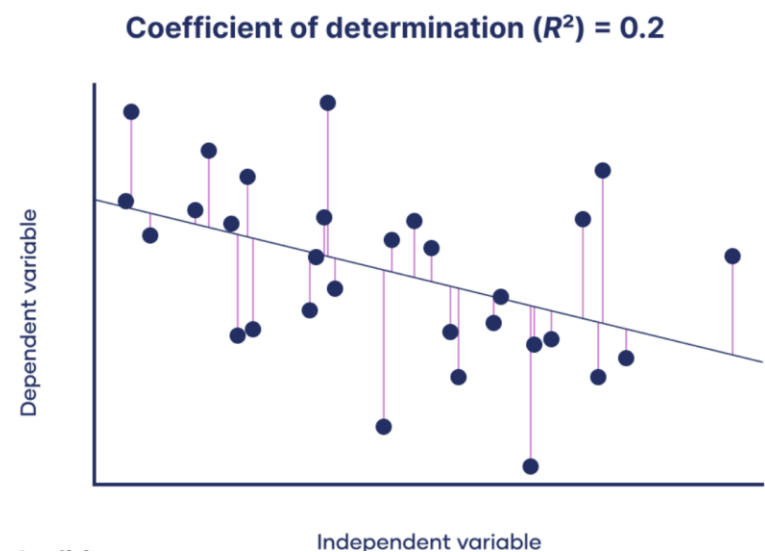
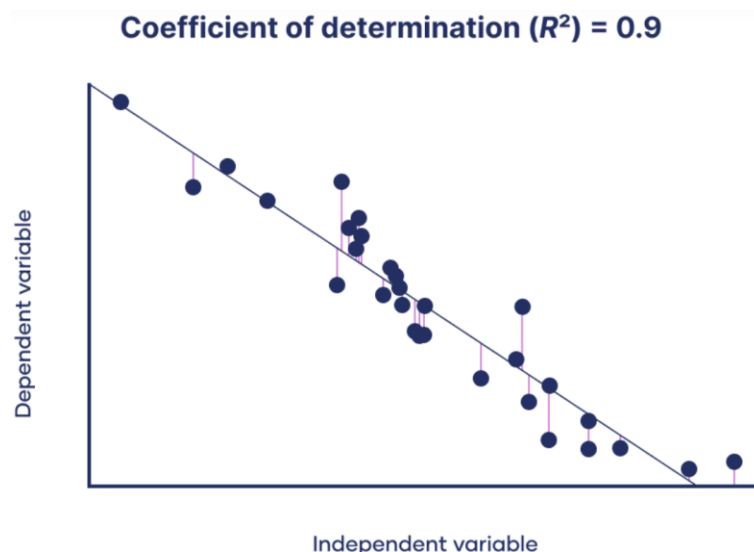
| | | | | | |
|------------|--------|--------|--------|--------|--------|
| | Test | Train | Train | Train | Train |
| | Train | Test | Train | Train | Train |
| | Train | Train | Test | Train | Train |
| | Train | Train | Train | Test | Train |
| | Train | Train | Train | Train | Test |
| Errors | 120.55 | 122.11 | 125.91 | 123.41 | 122.81 |
| Mean Error | 122.96 | | | | |

Error Metrics

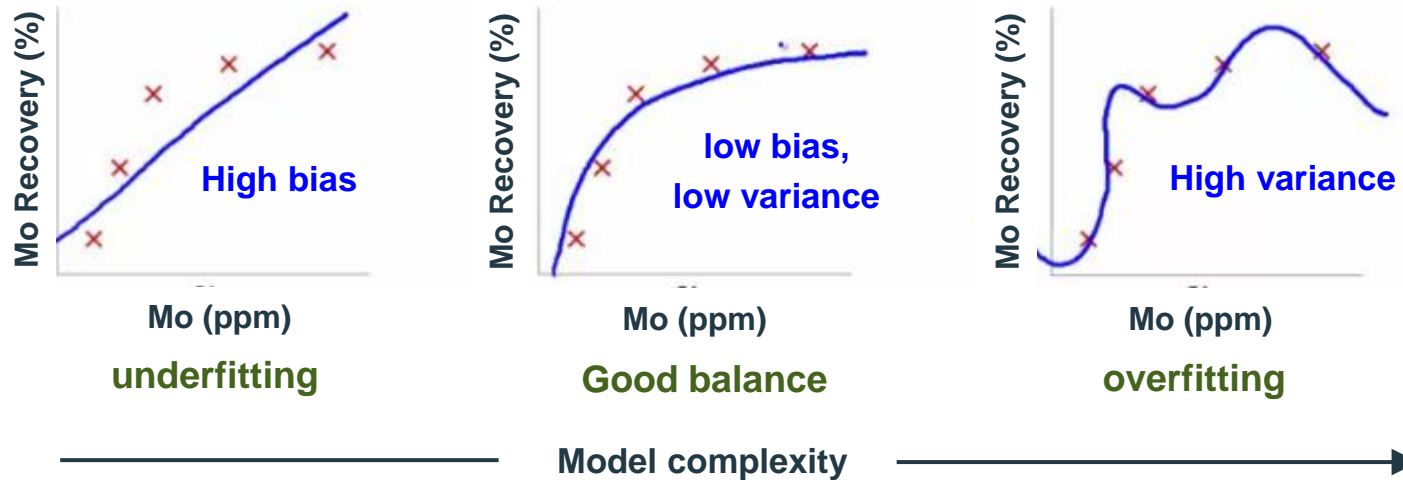
RMSE is the most used metric in regression.

$$RMSE = \sqrt{\frac{\sum_{i=1}^N (Predicted_i - Actual_i)^2}{N}}$$

Coefficient of Determination (r^2) determines the proportion of variance in the dependent variable that can be explained by the independent variable.

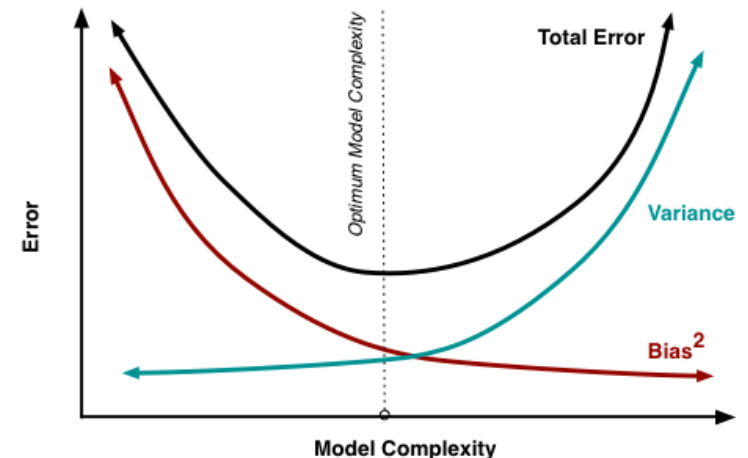


Bias-variance Tradeoff

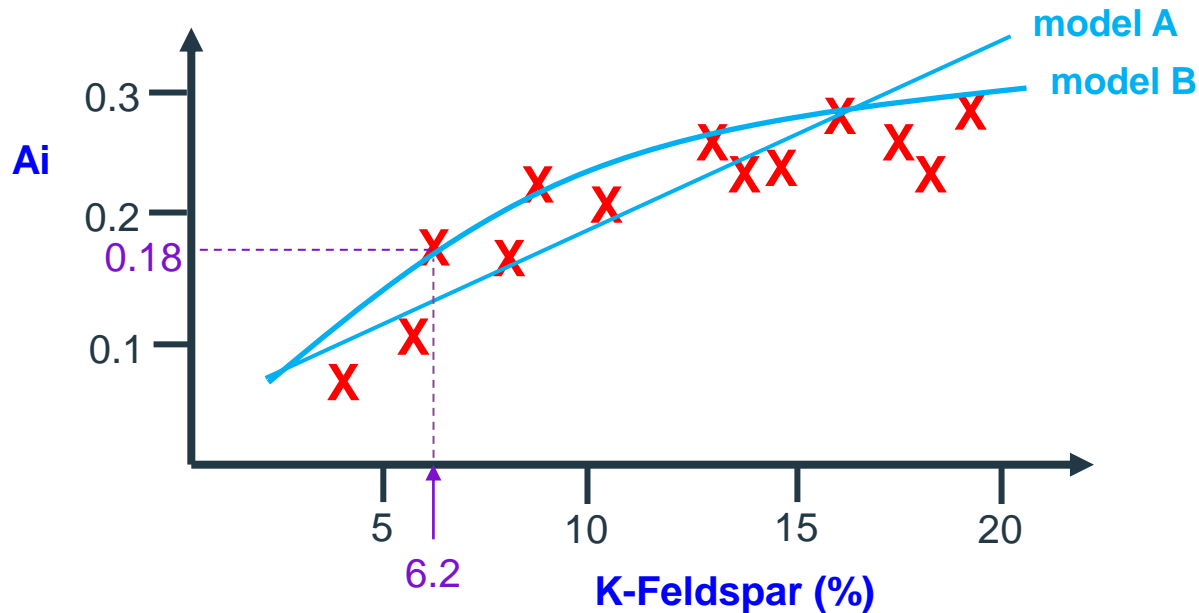


bias: difference between the model prediction and the actual value. Leads to high error on training and test data.

variance: variability of model prediction. Perform well on training data but has high error on test data.



How does a regression algorithm learn?



| predictor K-Feldspar (%) | target Ai |
|-----------------------------|--------------|
| 10.4 | 0.21 |
| 6.2 | 0.18 |
| 15.9 | 0.28 |
| ... | ... |

Regression: predict an infinite number of possible outputs. Model:

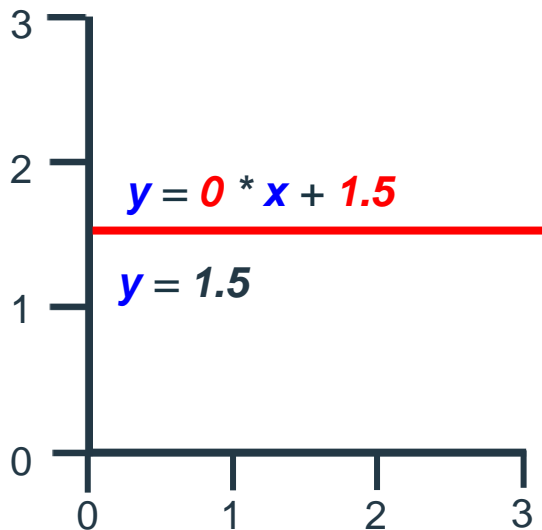
$$y = wx + b \rightarrow Ai = w * [K-Feldspar (%)] + b$$

Our goal is to find an algorithm that selects the most appropriate line/curve to fit the data. Which model is better, **model A** or **model B**? How can we choose the best?

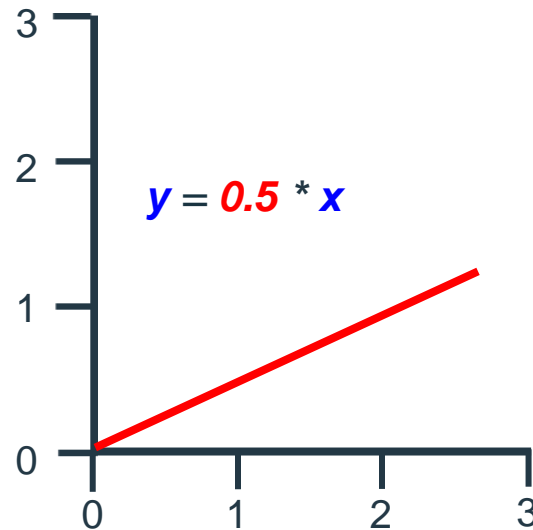
How does a regression algorithm learn?

Model: $y = wx + b$ | w, b = parameters (coefficients)

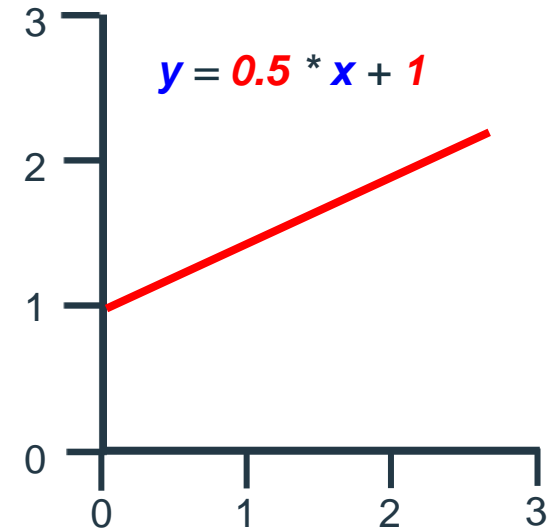
What w and b do?



$$w = 0$$
$$b = 1.5$$



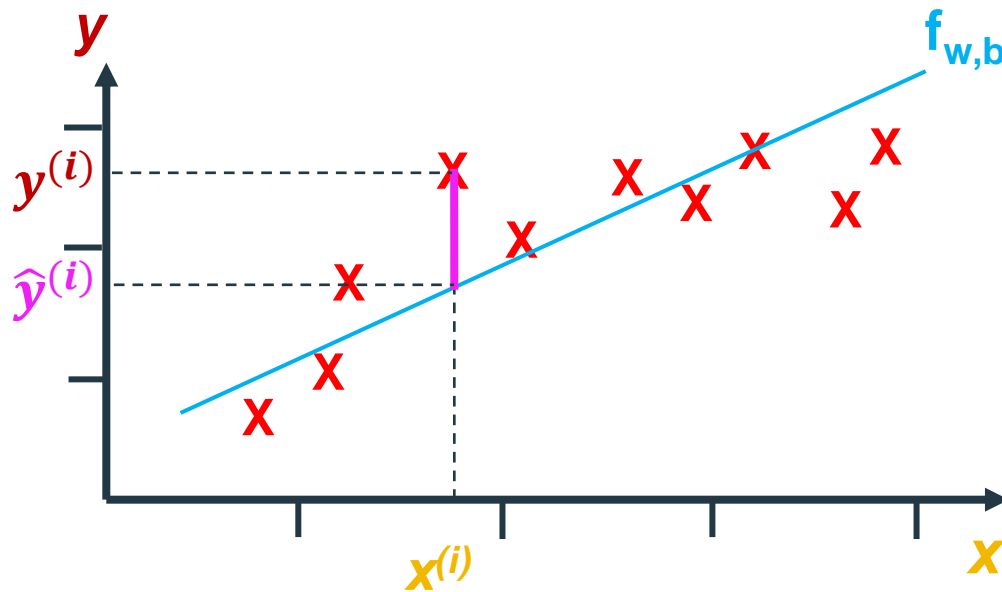
$$w = 0.5$$
$$b = 0$$



$$w = 0.5$$
$$b = 1$$

Modified from Andrew Ng, "machine learning specialization" (Coursera)

How does a regression algorithm learn?



prediction:

$$\hat{y}^{(i)} = f_{w,b}(x^{(i)})$$

cost function:

$$J(w,b) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2$$

Find w, b :

$\hat{y}^{(i)}$ is close to $y^{(i)}$ for all $(x^{(i)}, y^{(i)})$

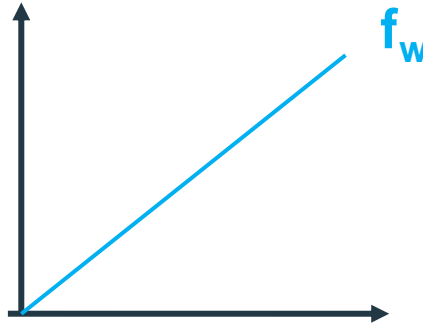
How does a regression algorithm learn?

Simplified Model:

$$f_{w,b}(x) = wx + b$$

if $b = 0$,

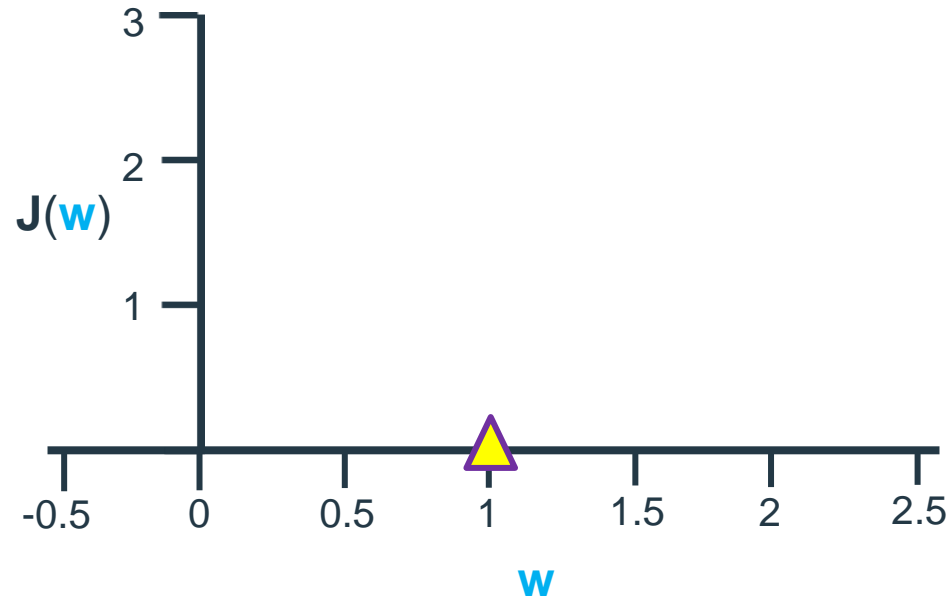
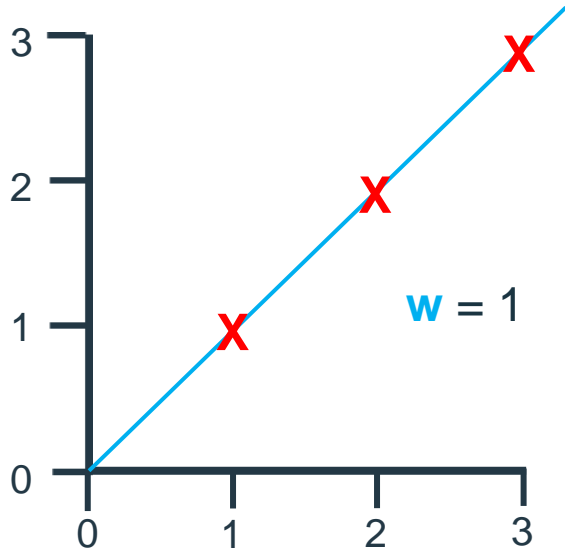
$$f_w(x) = wx$$



cost function:
$$J(w) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2$$

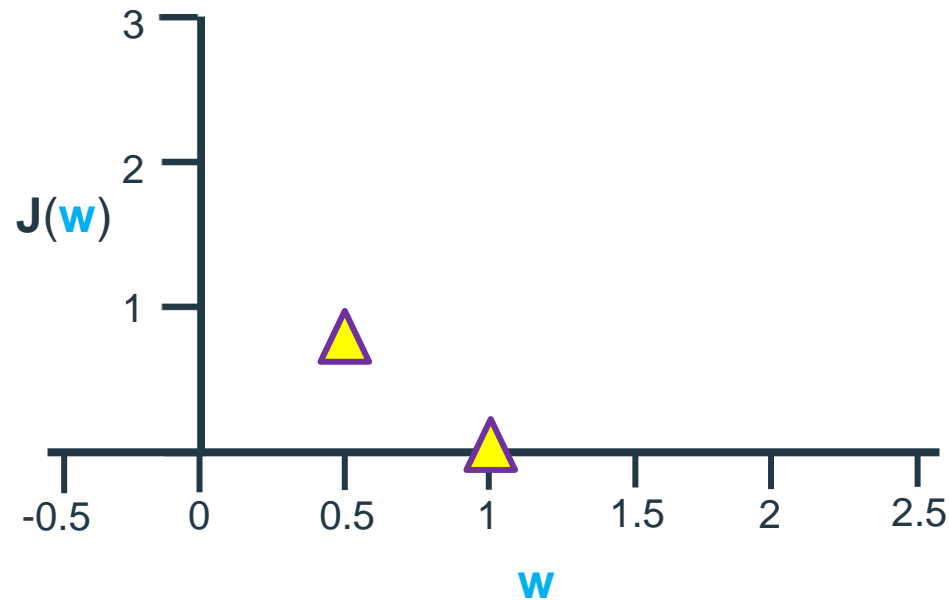
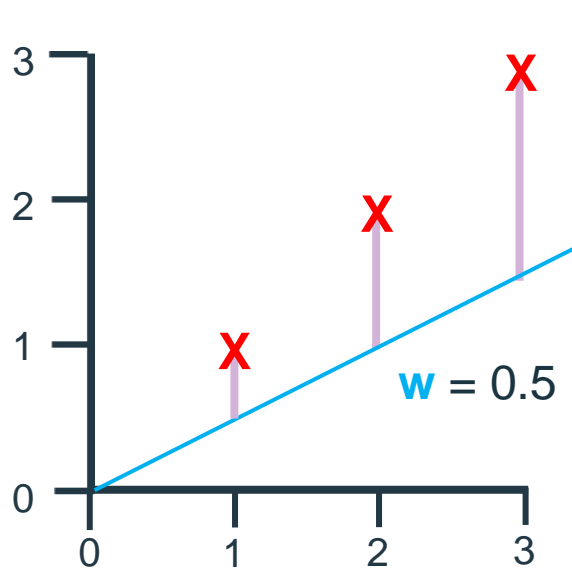
goal: minimize $J(w)$

How does a regression algorithm learn?



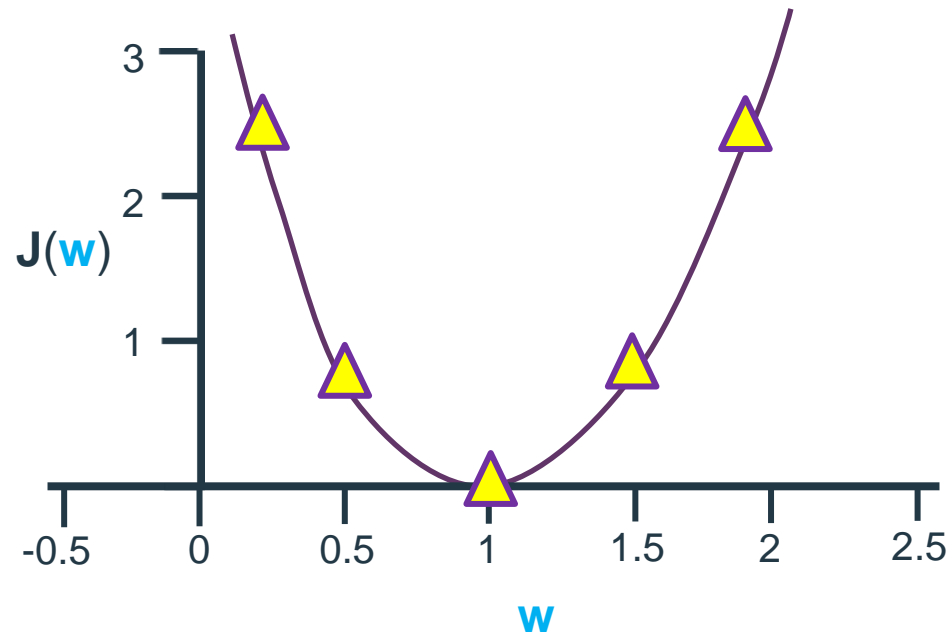
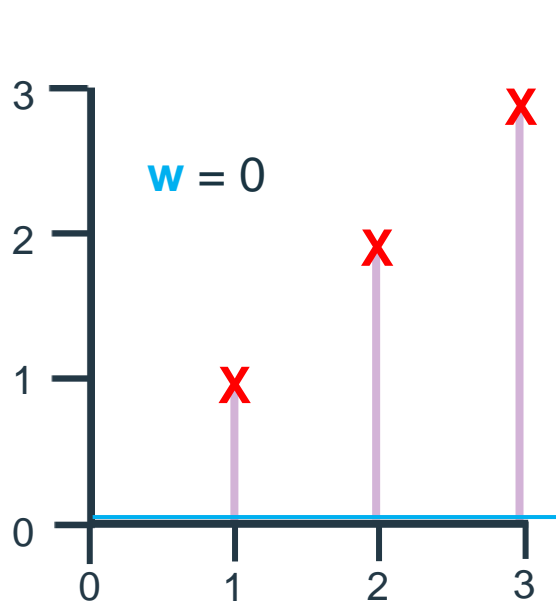
$$J(w = 1) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2 = \frac{1}{2m} (0^2 + 0^2 + 0^2) = 0$$

How does a regression algorithm learn?



$$J(w = 0.5) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2 = \frac{1}{2m} [(0.5-1)^2 + (2-1)^2 + (1.5-3)^2] = \frac{1}{2 \cdot 3} [3.5] = 0.58$$

How does a regression algorithm learn?



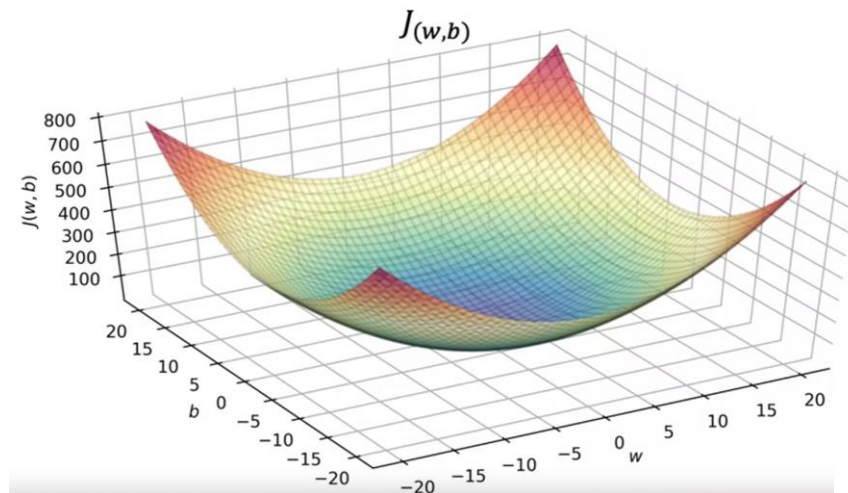
$$J(w = 0) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2 = \frac{1}{2m} (1^2 + 2^2 + 3^2) = \frac{1}{2 \cdot 3} [14] = 2.3$$

goal: minimize $J(w)$

How does a regression algorithm learn?

Any point represents a particular choice of w and b . The high in that point is the value of $J(w,b)$.

(GPT-3 parameters)



non-linear case

